

Technology Series



NoSQL

00




Presented by
Subhash EP



Subhash



Not
Only **SQL**



NoSQL Systems

Motivation

NoSQL: The Name

- “SQL” = Traditional relational DBMS
- Recognition over past decade or so:
Not every data management/analysis problem is best
solved using a traditional relational DBMS
- “NoSQL” = “No SQL” =
Not using traditional relational DBMS
- “No SQL” \neq Don't use SQL language

NoSQL: The Name

- “SQL” = Traditional relational DBMS
- Recognition over past decade or so:
Not every data management/analysis problem is best solved using a traditional relational DBMS
- “NoSQL” ^{exclusively} = “No SQL” =
Not using traditional relational DBMS
- “No SQL” \neq Don't use SQL language
- * “NoSQL” = ~~“Not Only SQL”~~

Not every data management/analysis problem is best solved using a traditional DBMS

Database Management System (DBMS) provides....

... efficient, reliable, convenient, and safe multi-user storage of and access to massive amounts of persistent data.

NoSQL Systems

Alternative to traditional relational DBMS

- + Flexible schema
- + Quicker/cheaper to set up
- + Massive scalability
- + Relaxed consistency → higher performance & availability
- No declarative query language → more programming
- Relaxed consistency → fewer guarantees

NoSQL Systems

Several incarnations

- MapReduce framework
- Key-value stores
- Document stores
- Graph database systems

MapReduce Framework

Originally from Google, open source Hadoop

- No data model, data stored in files
- User provides specific functions
map() reduce()
- System provides data processing “glue”, fault-tolerance,
scalability

Map and Reduce Functions

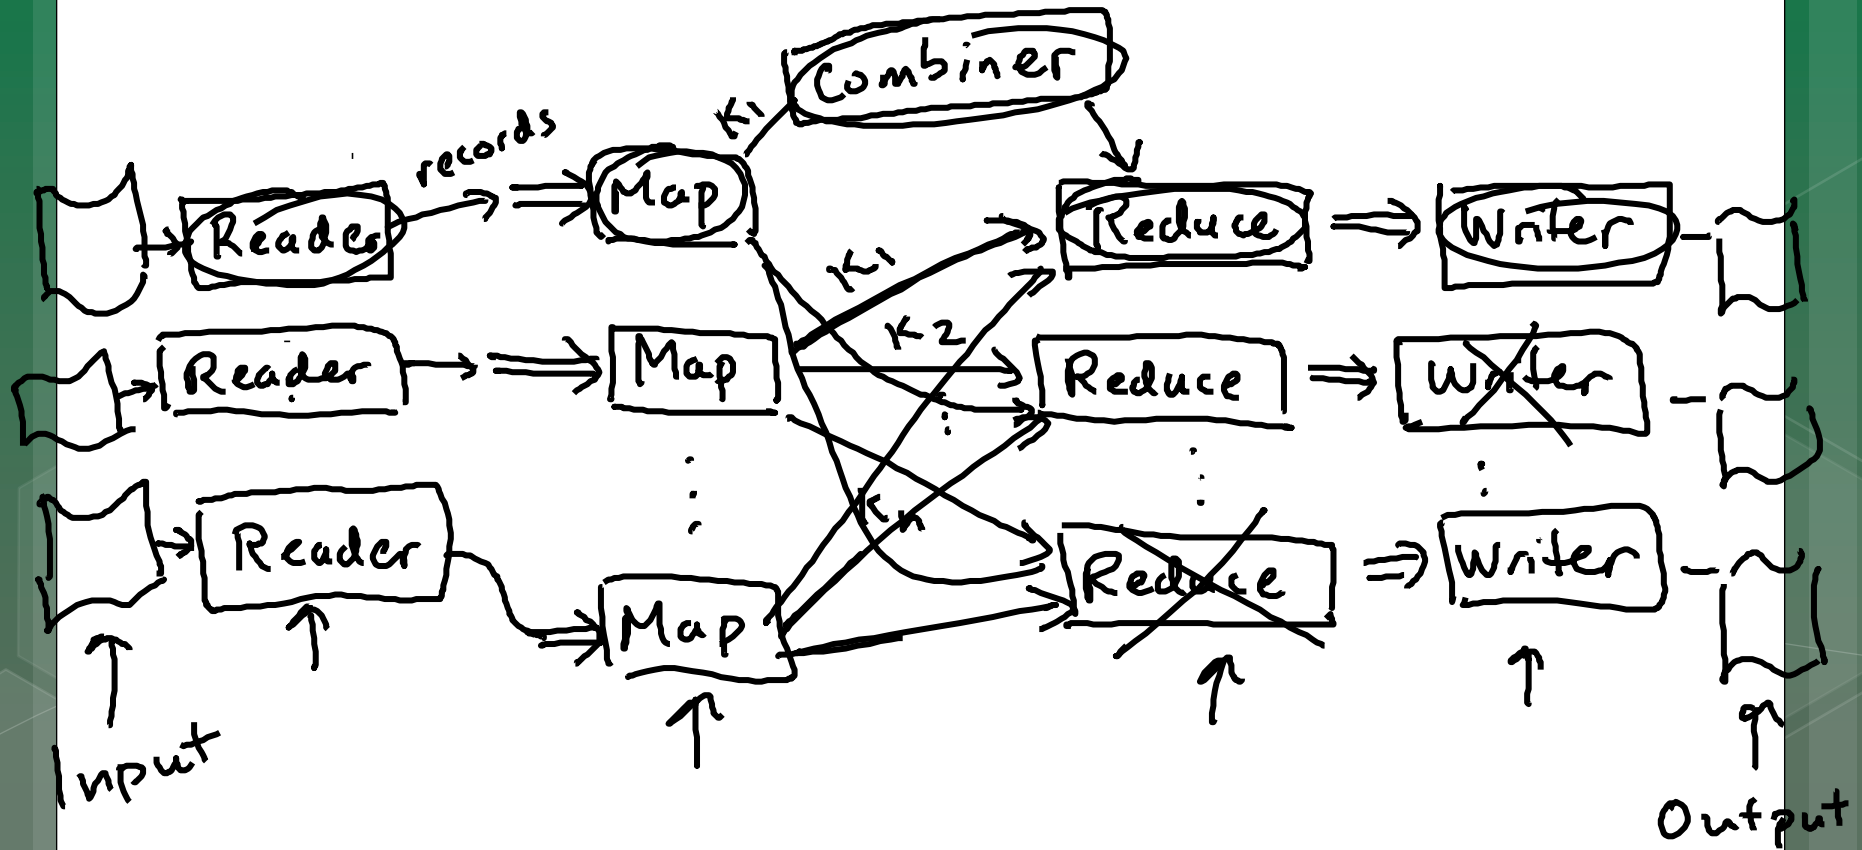
Map: Divide problem into subproblems

$\text{map}(\text{item}) \rightarrow 0 \text{ or more } \langle \text{key}, \text{value} \rangle \text{ pairs}$

Reduce: Do work on subproblems, combine results

$\text{reduce}(\text{key}, \text{list-of-values}) \rightarrow 0 \text{ or more records}$

MapReduce Architecture



MapReduce Example: Web log analysis

Each record: UserID, URL, timestamp, additional-info

Task: Count number of accesses for each domain (inside URL)

MapReduce Example (modified #1)

Each record: UserID, URL, timestamp, additional-info

Task: Total "value" of accesses for each domain based on additional-info

$\text{map}(\text{record}) \rightarrow \langle \text{domain}, \text{score} \rangle$
 $\text{reduce}(\text{domain}, \text{list of scores}) \rightarrow \langle \text{domain}, \text{sum} \rangle$

MapReduce Framework

- No data model, data stored in files
- User provides specific functions
- System provides data processing “glue”, fault-tolerance, scalability

MapReduce Framework

Schemas and declarative queries are missed

Hive – schemas, SQL-like query language

Pig – more imperative but with relational operators

- Both compile to “workflow” of Hadoop (MapReduce) jobs

Key-Value Stores

Extremely simple interface

- Data model: (key, value) pairs
- Operations: Insert(key,value), Fetch(key), Update(key), Delete(key)

Implementation: efficiency, scalability, fault-tolerance

- Records distributed to nodes based on key
- Replication
- Single-record transactions, “eventual consistency”

Key-Value Stores

Extremely simple interface

- Data model: (key, value) pairs
- Operations: Insert(key,value), Fetch(key), Update(key), Delete(key)
- Some allow (non-uniform) columns within value
- Some allow Fetch on range of keys

Example systems

- Google BigTable, Amazon Dynamo, Cassandra, Voldemort, HBase, ...

Document Stores

Like Key-Value Stores except value is document

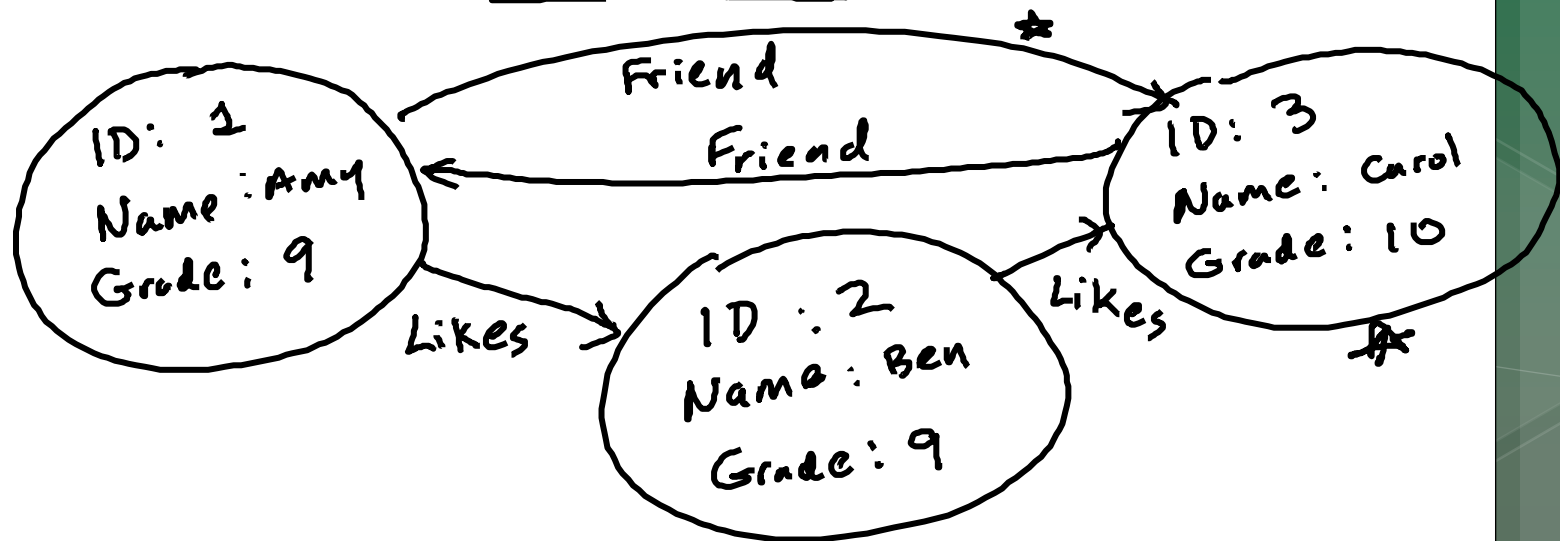
- **Data model:** (key, document) pairs
- **Document:** JSON, XML, other semistructured formats
- **Basic operations:** Insert(key,document),
Fetch(key),
Update(key), Delete(key)
- Also Fetch based on document contents

Example systems

- CouchDB, **MongoDB**, SimpleDB, ...

Graph Database Systems

- Data model: nodes and edges
- Nodes may have properties (including ID)
- Edges may have labels or roles



Graph Database Systems

- Interfaces and query languages vary
- Single-step versus “path expressions” versus full recursion
- Example systems
Neo4j, FlockDB, Pregel, ...

That's all



End of Session



subhash.ep@gmail.com
[linkedin.com/in/subhashep](https://www.linkedin.com/in/subhashep)

