



Technology Series

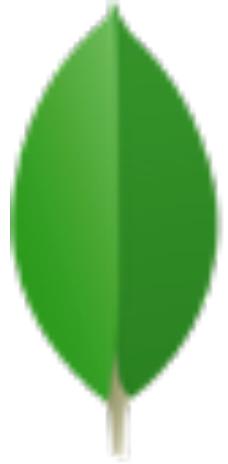
MongoDB

07

Presented by
Subhash EP



Subhash



mongoDB

{ name: mongo, type: DB }

What datatype is this?

	BIKE	COUNTER	LOVEMONGO	LOVESSQL	NAME	TS	WIFE
1	Felt	1	true	true	Marc	{ts '2010-10-20 18:12:21'}	Heather
2	Trek	2	true	false	Steve	{ts '2010-10-20 18:12:21'}	Paige
3	Giant	3	false	false	Dan	{ts '2010-10-20 18:12:21'}	Katie

And This?

BIKE	Felt
COUNTER	1
LOVESMONGO	true
LOVESSQL	true
NAME	Marc
TS	{ts '2010-10-20 21:03:45'}
WIFE	Heather

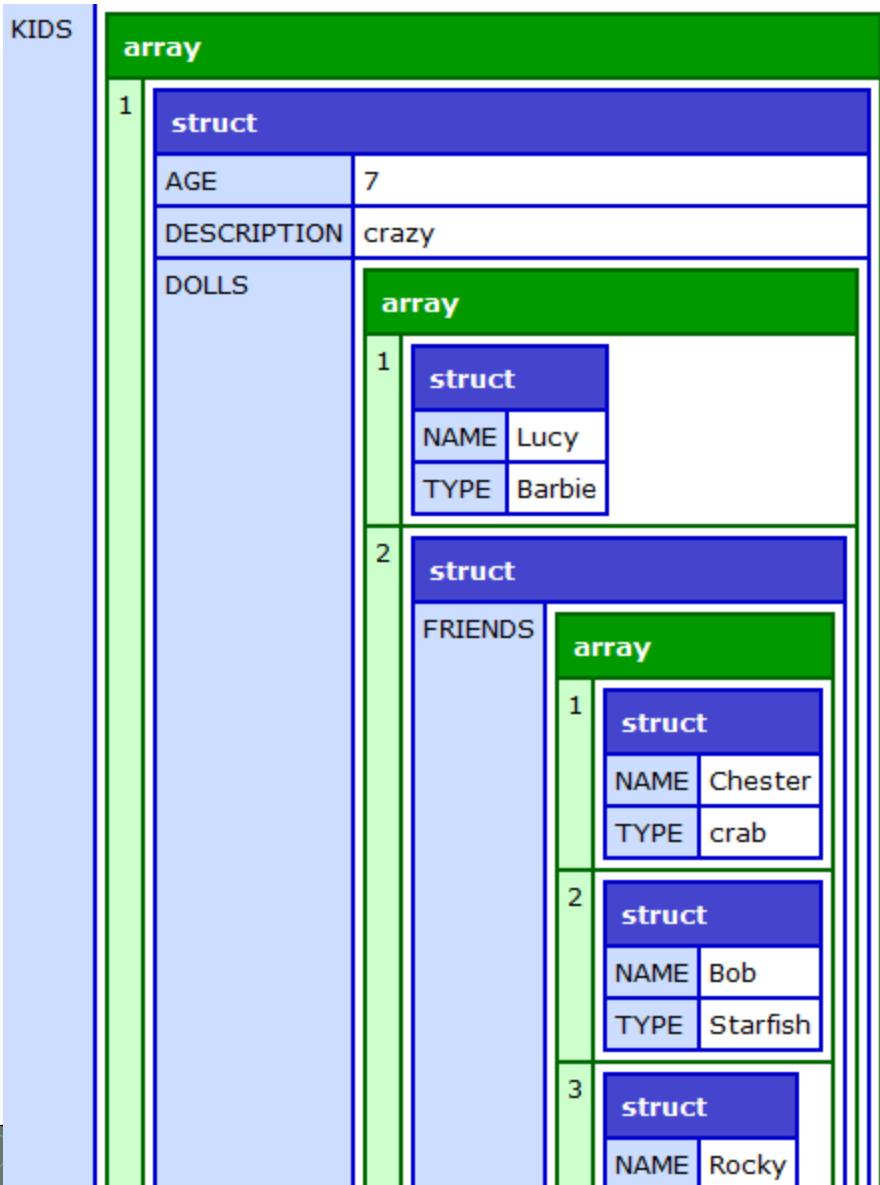
And This?

1	struct
BIKE	Specialized
COUNTER	1
KIDS	
NAME	Cool Dude 1
TS	{ts '2010-10-20 18:20:01'}
WIFE	Smokin hot wife 1

2	struct
BIKE	Specialized
COUNTER	2
KIDS	
NAME	Cool Dude 2
TS	{ts '2010-10-20 18:20:01'}
WIFE	Smokin hot wife 2

3	struct
BIKE	Specialized
COUNTER	3
KIDS	
NAME	Cool Dude 3

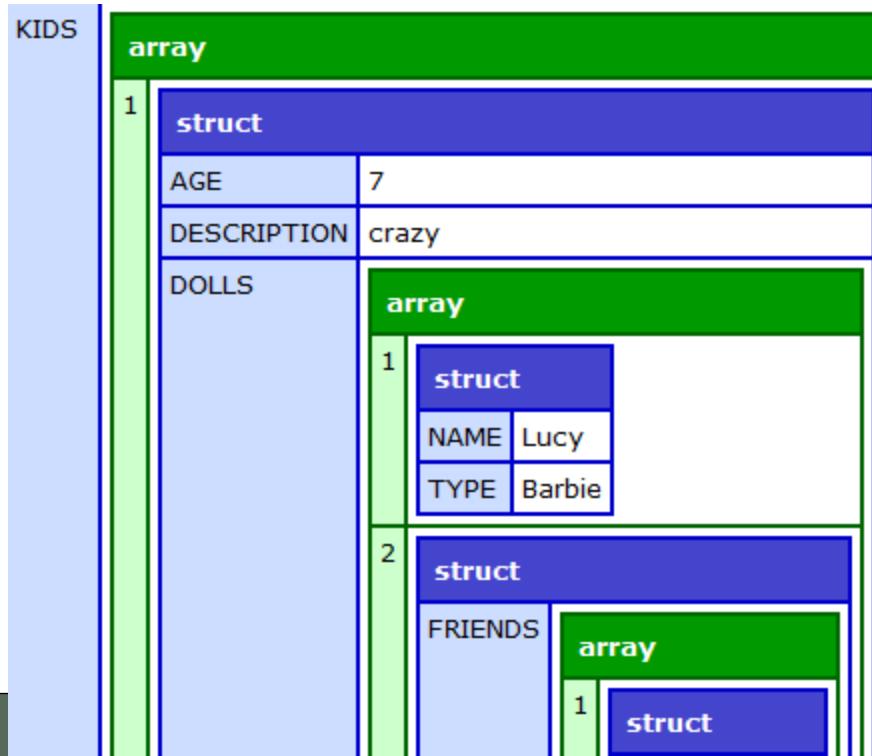
And This?



And the first picture is different from that last one how?

	BIKE	COUNTER	LOVEMONGO	LOVESSQL	NAME	TS	WIFE
1	Felt	1	true	true	Marc	{ts '2010-10-20 18:12:21'}	Heather
2	Trek	2	true	false	Steve	{ts '2010-10-20 18:12:21'}	Paige
3	Giant	3	false	false	Dan	{ts '2010-10-20 18:12:21'}	Katie

VS



Mindshift

From Queries

To Arrays of Structures

MongoDB is...

- A Schema-less, Document-Oriented datastore
- Documents are roughly equivalent to KV structs:
 - Keys, values
 - Values can be other documents (i.e. embedded)
- Documents are **searchable**

MongoDB's Data Model

- A Database has “Collections”
 - Collections have “Documents”
 - Documents have “Fields”
 - Fields are key = value pairs
- A Collection does not enforce the structure of its documents*

*i.e. Schemaless

Install and Run MongoDB

Download from mongodb.org

Unzip

Create data directory

```
>mkdir c:\data\db
```

Run MongoDB (**mongod**):

```
>cd c:\mongodb-1.6.3\bin
```

```
>mongod
```

Run Mongo shell (**mongo**):

```
>mongo
```

Yes, that's it

The Mongo Shell

>mongo

>help()

>show dbs

>use <dbname>

>show collections

>db.collectionName.findOne()

>db.collectionName.find()

>db.help()

>db.collectionName.help()

Insert

```
>
>
> db.people.insert( {NAME: "Steve", BIKE: "Trek", WIFE: "Paige", TS: new Date()} )
>
> db.people.findOne( {NAME: "Steve"} )
{
  "_id" : ObjectId("4cbcef22d710000000051cc"),
  "NAME" : "Steve",
  "BIKE" : "Trek",
  "WIFE" : "Paige",
  "TS" : "Mon Oct 18 2010 21:09:06 GMT-0400 (Eastern Daylight Time)"
}
>
```

Update

```
> db.people.update( {NAME: "Steve"}, { $set: { DOG: "Apollo" } } )
>
> db.people.findOne( {NAME: "Steve"} )
{
  "BIKE" : "Trek",
  "DOG" : "Apollo",
  "NAME" : "Steve",
  "TS" : "Mon Oct 18 2010 21:09:06 GMT-0400 (Eastern Daylight Time)",
  "WIFE" : "Paige",
  "_id" : ObjectId("4cbcef22d71000000051cc")
}
>

> db.people.update( {NAME: "Steve"}, { $unset: { DOG: true } } )
>
> db.people.findOne( {NAME: "Steve"} )
{
  "BIKE" : "Trek",
  "NAME" : "Steve",
  "TS" : "Mon Oct 18 2010 21:09:06 GMT-0400 (Eastern Daylight Time)",
  "WIFE" : "Paige",
  "_id" : ObjectId("4cbcef22d71000000051cc")
}
```

Updating Components

```
> db.Components.find();
{ "_id" : 1, "SerialNumber" : "1-002" }
> db.Components.update({"_id":1}, {"$set":{"ModelNumber":"Model1"}});
> db.Components.find();
{ "ModelNumber" : "Model1", "SerialNumber" : "1-002", "_id" : 1 }
>
```

- \$set keyword used for partial updates
- Without \$set keyword entire document is replaced
- {multi : true} to update multiple documents

\$set, \$unset, \$inc, \$push, \$pushAll, \$pull, \$pullAll, \$pop,
\$addToSet, \$rename, \$bit, \$ positional operator

Remove

```
> db.people.remove( {NAME: "Steve"} )
> db.people.findOne( {NAME: "Steve"} )
null
>
```

Deleting Components

`db.Components.remove();` Works like `.find()`

`db.Components.drop();` Drops collection

`db.dropDatabase();` Drops database

Finding by nested fields

```
> db.people.insert( {NAME: "Bob", KIDS: [ {NAME: "Sally", AGE: 3}, {NAME: "Billy", AGE: 7} ] })
```

```
> db.people.find({NAME: "Bob"}).pretty()
```

Finding by nested fields

```
> db.people.insert( {NAME: "Bob", KIDS: [ {NAME: "Sally", AGE: 3}, {NAME: "Billy", AGE: 7} ] } )
>
> db.people.findOne({NAME: "Bob"})
...
{
  "_id" : ObjectId("4cbcf1632d710000000051cf"),
  "NAME" : "Bob",
  "KIDS" : [
    {
      "NAME" : "Sally",
      "AGE" : 3
    },
    {
      "NAME" : "Billy",
      "AGE" : 7
    }
  ]
}
```

Finding by nested fields

```
> db.people.insert( {NAME: "Bob", KIDS: [ {NAME: "Sally", AGE: 3}, {NAME: "Billy", AGE: 7} ] } )
> db.people.findOne({ "KIDS.NAME": "Billy"})
{
  "_id" : ObjectId("4cbcfc1632d7100000000051cf"),
  "NAME" : "Bob",
  "KIDS" : [
    {
      "NAME" : "Sally",
      "AGE" : 3
    },
    {
      "NAME" : "Billy",
      "AGE" : 7
    }
  ]
}
>
```

Finding with Conditionals

```
> db.people.insert( {NAME: "Shaggy", KIDS: [ {NAME: "Scrappy", AGE: 3}, {NAME: "Scooby", AGE: 5} ] } )
> db.people.insert( {NAME: "Fred", KIDS: [ {NAME: "Biff", AGE: 13}, {NAME: "Muffy", AGE: 10} ] } )
> db.people.insert( {NAME: "Thelma", KIDS: [ {NAME: "Simon", AGE: 4}, {NAME: "Stella", AGE: 2} ] } )
>
```

Finding with Conditionals

```
> db.people.insert( {NAME: "Shaggy", KIDS: [ {NAME: "Scrappy", AGE: 3}, {NAME: "Scooby", AGE: 5} ] } )
> db.people.insert( {NAME: "Fred", KIDS: [ {NAME: "Biff", AGE: 13}, {NAME: "Muffy", AGE: 10} ] } )
> db.people.insert( {NAME: "Thelma", KIDS: [ {NAME: "Simon", AGE: 4}, {NAME: "Stella", AGE: 2} ] } )
>
> db.people.find({ "KIDS.AGE" : { $in: [3,4] } })
{ "_id" : ObjectId("4cbcfc3f92d710000000051d3"), "NAME" : "Shaggy", "KIDS" :
  [
    {
      "NAME" : "Scrappy",
      "AGE" : 3
    },
    {
      "NAME" : "Scooby",
      "AGE" : 5
    }
  }
{ "_id" : ObjectId("4cbcfc4012d710000000051d5"), "NAME" : "Thelma", "KIDS" :
  [
    {
      "NAME" : "Simon",
      "AGE" : 4
    },
    {
      "NAME" : "Stella",
      "AGE" : 2
    }
  ]
}
```

What happens if I search on a field a Document doesn't have?

```
> db.people.insert( {NAME: "Daphne", SWINGER: true })  
> db.people.find({ "KIDS.AGE": { $in: [3,4] } })
```

What happens if I search on a field a Document doesn't have?

```
> db.people.insert( {NAME: "Daphne", SWINGER: true })  
>  
> db.people.find({ "KIDS.AGE" : { $in: [3,4] } }, {})  
{ "_id" : ObjectId("4cbcfc3f92d710000000051d3"), "NAME" : "Shaggy", "KIDS" : [  
    {  
        "NAME" : "Scrappy",  
        "AGE" : 3  
    },  
    {  
        "NAME" : "Scooby",  
        "AGE" : 5  
    }  
], }  
{ "_id" : ObjectId("4cbcfc4012d710000000051d5"), "NAME" : "Thelma", "KIDS" : [  
    {  
        "NAME" : "Simon",  
        "AGE" : 4  
    },  
    {  
        "NAME" : "Stella",  
        "AGE" : 2  
    }  
], }  
>
```

What happens if I search on a field a Document doesn't have?

```
>  
>  
-> db.people.find({SWINGER: true})  
{ "_id" : ObjectId("4cbcfc5392d710000000051d6") , "NAME" : "Daphne", "SWINGER" : true }  
>
```

Can I Find Documents without XXX?

```
> db.people.find({KIDS: {$exists: false} })  
{ "_id" : ObjectId("4cce10cbff350000000034a2"), "NAME" : "Daphne", "SWINGER" : true }
```

Sorting? Limiting? Skipping?

Have you ever had to “Page” in a web app?



```
> db.people.insert( {NAME: "Thelma", AGE: 25 } )
> db.people.insert( {NAME: "Daphne", AGE: 33 } )
> db.people.insert( {NAME: "Shaggy", AGE: 19 } )
> db.people.insert( {NAME: "Scrappy", AGE: 3 } )
>
```

Sorting?

```
> db.people.insert( {NAME: "Thelma", AGE: 25 } )
> db.people.insert( {NAME: "Daphne", AGE: 33 } )
> db.people.insert( {NAME: "Shaggy", AGE: 19 } )
> db.people.insert( {NAME: "Scrappy", AGE: 3 } )
>
>
> db.people.find().sort({AGE: -1})
[{"_id": ObjectId("4cbcfc8042d710000000051d8"), "NAME": "Daphne", "AGE": 33},
 {"_id": ObjectId("4cbcfc7f12d710000000051d7"), "NAME": "Thelma", "AGE": 25},
 {"_id": ObjectId("4cbcfc8262d710000000051d9"), "NAME": "Shaggy", "AGE": 19},
 {"_id": ObjectId("4cbcfc82e2d710000000051da"), "NAME": "Scrappy", "AGE": 3}]
```

1 = true, 1 = asc, 1 = include

-1 = false, -1 = desc, -1 = exclude

(depends on context)

Sorting? Limiting?

```
> db.people.insert( {NAME: "Thelma", AGE: 25 } )
> db.people.insert( {NAME: "Daphne", AGE: 33 } )
> db.people.insert( {NAME: "Shaggy", AGE: 19 } )
> db.people.insert( {NAME: "Scrappy", AGE: 3 } )
>
>
> db.people.find().sort({AGE: -1})
{
  "_id" : ObjectId("4cbcfc8042d710000000051d8"),
  "NAME" : "Daphne",
  "AGE" : 33
},
{
  "_id" : ObjectId("4cbcfc7f12d710000000051d7"),
  "NAME" : "Thelma",
  "AGE" : 25
},
{
  "_id" : ObjectId("4cbcfc8262d710000000051d9"),
  "NAME" : "Shaggy",
  "AGE" : 19
},
{
  "_id" : ObjectId("4cbcfc82e2d710000000051da"),
  "NAME" : "Scrappy",
  "AGE" : 3
}
>
> db.people.find().sort({AGE: -1}).limit(2)
{
  "_id" : ObjectId("4cbcfc8042d710000000051d8"),
  "NAME" : "Daphne",
  "AGE" : 33
},
{
  "_id" : ObjectId("4cbcfc7f12d710000000051d7"),
  "NAME" : "Thelma",
  "AGE" : 25
}
>
```

sort(), skip(), limit()

Sorting + Limiting + Skipping =

```
> db.people.insert( {NAME: "Thelma", AGE: 25 } )
> db.people.insert( {NAME: "Daphne", AGE: 33 } )
> db.people.insert( {NAME: "Shaggy", AGE: 19 } )
> db.people.insert( {NAME: "Scrappy", AGE: 3 } )
>
>
> db.people.find().sort({AGE: -1})
{ "_id" : ObjectId("4cbcfc8042d710000000051d8"), "NAME" : "Daphne", "AGE" : 33 }
{ "_id" : ObjectId("4cbcfc7f12d710000000051d7"), "NAME" : "Thelma", "AGE" : 25 }
{ "_id" : ObjectId("4cbcfc8262d710000000051d9"), "NAME" : "Shaggy", "AGE" : 19 }
{ "_id" : ObjectId("4cbcfc82e2d710000000051da"), "NAME" : "Scrappy", "AGE" : 3 }
>
> db.people.find().sort({AGE: -1}).limit(2)
{ "_id" : ObjectId("4cbcfc8042d710000000051d8"), "NAME" : "Daphne", "AGE" : 33 }
{ "_id" : ObjectId("4cbcfc7f12d710000000051d7"), "NAME" : "Thelma", "AGE" : 25 }
>
>
> db.people.find().sort({AGE: -1}).skip(1).limit(2)
{ "_id" : ObjectId("4cbcfc7f12d710000000051d7"), "NAME" : "Thelma", "AGE" : 25 }
{ "_id" : ObjectId("4cbcfc8262d710000000051d9"), "NAME" : "Shaggy", "AGE" : 19 }
```

sort(), skip(), limit()

count() and distinct()

```
> db.people.insert({NAME: 'Shaggy', AGE: 25})  
> db.people.insert({NAME: 'Fred', AGE: 28})  
> db.people.insert({NAME: 'Daphne', AGE: 27})  
> db.people.insert({NAME: 'Thelma', AGE: 24})  
>  
>  
> db.people.count()  
10
```

count() and distinct()

```
> db.people.insert({NAME: 'Shaggy', AGE: 25})  
> db.people.insert({NAME: 'Fred', AGE: 28})  
> db.people.insert({NAME: 'Daphne', AGE: 27})  
> db.people.insert({NAME: 'Thelma', AGE: 24})  
>  
>  
> db.people.count()  
10  
>  
> db.people.count({AGE: {$gt: 24}})  
3  
>
```

count() and distinct()

```
> db.people.insert({NAME: 'Shaggy', AGE: 25})  
> db.people.insert({NAME: 'Fred', AGE: 28})  
> db.people.insert({NAME: 'Daphne', AGE: 27})  
> db.people.insert({NAME: 'Thelma', AGE: 24})  
>  
>  
> db.people.count()  
10  
>  
> db.people.count({AGE: {$gt: 24}})  
3  
>  
> db.people.distinct( 'AGE' )  
[ 24, 25, 27, 28 ]
```

count() and distinct()

```
> db.people.insert({NAME: 'Shaggy', AGE: 25})
> db.people.insert({NAME: 'Fred', AGE: 28})
> db.people.insert({NAME: 'Daphne', AGE: 27})
> db.people.insert({NAME: 'Thelma', AGE: 24})
>
>
> db.people.count()
10
>
> db.people.count({AGE: {$gt: 24}})
3
>
> db.people.distinct( 'AGE' )
[ 24, 25, 27, 28 ]
>
> db.people.distinct( 'NAME' )
[
    "Cool Dude 1",
    "Cool Dude 2",
    "Cool Dude 3",
    "Cool Dude 4",
    "Cool Dude 5",
    "Daphne",
    "Fred",
    "Marc",
    "Shaggy",
    "Thelma"
]
>
```

Schema Design Guideline

Embed when you can:

```
{ NAME = "Marc", KIDS = [{NAME="Lexie"}, {NAME="Sidney"}] }
```

Relate when you must:

```
{ NAME = "Marc",
  KIDS = [
    ObjectId(4cd0bcc754503766b7dcfd6a),
    ObjectId(4cd0bcc754503766b7dcfd7a)
  ]
}
```

Why MongoDB: The Official Word

MongoDB is **Big** and **Fast**

(just like **every other** NoSQL Datastore)

Why MongoDB: Performance

- No Joins + No multi-row transactions
 - = Fast Reads
 - = Fast Writes (b/c you write to fewer tables, no trans. log)
- Async writes
 - = you don't wait for inserts to complete
 - (optional, though)
- Secondary Indexes
 - = Index on embedded document fields for superfast ad-hoc queries
 - Indexes live in RAM

Why MongoDB: Scalability / Availability

- No Joins + No multi-row transactions
 - = Horizontally Scalable
 - = Built for Distribution of data and computation
 - Sharding for distributed data
 - MapReduce for distributed computation
- Built-in Replication
 - Easy to configure
 - Reads can be distributed across slaves
 - Automatic failover when using Replica Sets

Shard for scalability, **replicate** for availability

Why MongoDB: R.A.D.

- Documents (think: objects) with embedded documents feel more natural compared with DB / ORM
- Ad-hoc queries don't require "view" creation or other config... code-n-go
- No penalty for schema evolution (true for most NoSQL stores)
- Atomic Modifiers for safe in-place updates (no locking code around single documents)
- Built-in functions for Mongo-as-a-queue (I LOVE these!)

Why MongoDB: Friendly License

- You can use MongoDB in your corporate app
- You do not have to open source your corporate app
- From the docs:

“ To reiterate, you **only** need to provide the source for the MongoDB server and not your application”

<http://www.mongodb.org/display/DOCS/Licensing>

Why Not MongoDB? Vs. Itself

- Still kinda new
 - Though used in production at major sites
- Evolving rapidly
 - This is a good thing, though you'll want to keep up with new releases
- Not meant for 32Bit
 - You'll have a 2GB limit per mongod process if you do
- No single-server durability
 - If you're not running replication, and it crashes, you will lose data (they're working on this)
 - Just replicate, period

Awesome: Query Operators

\$ne:

```
db.people.find( { NAME: { $ne: "Shaggy" } } )
```

\$nin:

```
db.people.find( { NAME: { $nin: ["Shaggy", "Daphne"] } } )
```

\$gt and \$lt:

```
db.people.find( AGE: { {$gt: 30, $lt: 35} } )
```

\$size (eg, people with exactly 3 kids):

```
db.people.find( KIDS: { $size: 3 } )
```

regex: (eg, names starting with Ma or Mi)

```
db.people.find( { NAME: / ^M(a|i) / } )
```

Awesome: Atomic Modifiers

\$inc

```
db.pageviews.update( {URL: 'http://myurl.com'}, {$inc: {N: 1}} )
```

\$set

```
db.people.update( {NAME: 'Steve'}, {$set: {Age: 35}} )
```

\$push (for atomically adding values to an array)

```
db.people.update( {NAME: 'Steve'}, {$push: {KIDS: {NAME: 'Sylvia', AGE: 3}}})
```

findAndModify()

```
db.tasks.findAndModify(  
  query: {STATUS: 'pending'},  
  sort: {PRIORITY: -1},  
  update: {$set: {STATUS: 'running', TS: new Date()}}  
)
```

Awesome: ObjectId

- BSON (binary-encoded JSON)
- Unique: combination of
 - Time
 - Machine
 - Process id
 - Incrementer
- Mongo will automatically add this `_id` field

That's all



End of Session

©

subhash.ep@gmail.com
linkedin.com/in/subhashep